

YOWON AI

Autonomous AI Jury Platform

Yowon AI
Research Project

73	CONDITIONAL_APPROVE	MEDIUM
Overall Score	Deployment Decision	Risk Level

Generated: June 18, 2026 06:13 UTC

Executive Summary

Overview

The project has an overall score of 73 with a `CONDITIONAL_APPROVE` verdict and `MEDIUM` risk level. It demonstrates strengths in technical implementation, security considerations, and scalability, but falls short in innovation and impact.

Evaluation Context

Item	Value
Selected Project Type	Research Project
AI Detected Type	Hackathon Project (60% confidence)
Scoring Rubric Used	Research Project
Evaluation Standard	Research novelty, academic contribution, experimental rigor, reproducibility, baseline benchmarking, and publication potential
Score Band	Strong
Confidence	87/100
Evidence Quality	Exceptional
Repository Completeness	100/100


Repository Statistics

Item	Value
Total Files	128
Code Files	105
Documentation Files	9
Presentation Files	0
Test Files	14
Configuration Files	6
Deployment Files	0
Data Files	1
Source Modules	108
Meaningful Files	119
Repository Completeness Score	100

Score Table


Category	Score	Grade
Forge	76/100	B
Sentinel	74/100	B
Scale	74/100	B
Visionary	72/100	B
Guardian	69/100	B

Category Score Chart



Radar Chart

Radar profile of calibrated specialist scores



Risk Visualization

Risk Area	Signal	Risk Level
Security Risk	MEDIUM	
Evidence Gaps	3	
Deployment Readiness	73/100	

Strengths

- Modular architecture
- Real-time SSE streaming
- Integration of RAG

Weaknesses

- Innovation
- Impact

Positive Factors

- Source code detected
- Documentation present
- Automated testing detected
- Machine learning implementation
- Dataset artifacts detected
- Backend architecture present
- Modular architecture detected
- Machine learning model detected

Missing Evidence

- No deployment evidence
- No baseline comparison (-10)
- No citation or literature-grounding evidence (-8)

Deployment Roadmap

1. Refactor code for better modularity and maintainability
2. Implement additional security measures to improve security score

Architecture Summary

Detected frontend, backend, database, ml pipeline, api layer, authentication, integrations, queues, vector databases, agent systems

Detected Technologies

- Celery
- ChromaDB
- CrewAI
- FastAPI
- LangChain
- React
- SQLAlchemy
- scikit-learn

Detected Algorithms

- Agent System
- RAG
- Random Forest
- Recommendation System
- Search/Ranking

Evidence Found

- Machine learning model
- Custom algorithm
- REST API
- Database
- Authentication
- Testing
- Security implementation
- External integrations
- Queue/background jobs
- Vector database
- Agent system

Evidence Missing

- Docker/deployment

Project Type Justification

Detected Hackathon Project from code/docs/repository signals.

Calibration Explanation

Implementation evidence increased confidence: Machine learning model, Custom algorithm, REST API, Database, Authentication, Testing. Missing evidence primarily reduces confidence unless required by project type: Docker/deployment. Community impact contributed a bounded signal (6/100, capped at 15% of final score). 9 calibration adjustment(s) were applied.

Confidence Sources

- Repository completeness: 100/100
- Documentation quality: present
- Evidence coverage: 100/100
- Agent agreement: inferred from calibrated specialist scores
- Dataset artifact detected

Penalties

- technical: Dataset artifacts strengthen technical evidence (+4)
- technical: REST/API implementation detected (+4)
- technical: Database integration detected (+3)
- technical: External integration evidence detected (+3)
- technical: No baseline comparison (-10)
- security: Authentication implementation detected (+4)
- scalability: Advanced architecture component detected (+5)
- innovation: Custom algorithmic implementation detected (+6)
- innovation: No citation or literature-grounding evidence (-8)

Detailed Findings

Forge Analysis

Forge Analysis

Technical Score:

76/100

Strengths:

Modular architecture

Real-time SSE streaming

Weaknesses:

Missing pitch deck or PDF

No deployment files

Risks:

Lack of commercial deployment context

Limited code volume for production

Sentinel Analysis

Sentinel Analysis

Security Score:

74/100

Risk Level:

LOW

Findings:

None evidenced.

Recommendations:

Run a security review and document controls.

Confidence:

90%

Visionary Assessment

Visionary Analysis

Innovation Score:

72/100

Scalability Score:

74/100

Differentiators:

Integration of RAG

Custom algorithms

Risks:

Needs clearer deployment strategy

Recommendations:

Document novelty, baseline comparison, and scale assumptions.

Guardian Impact Analysis

Guardian Analysis

Impact Score:

69/100

Top Risks:

Lack of deployment experience

No production testing

Insufficient security measures

Expected Impact:

Model drift in agents

Data quality issues

API latency

Top Failure Predictions

- 1. Model drift in agents
- 2. Data quality issues
- 3. API latency

Scalability Assessment

Visionary Scalability Assessment

Scalability Score:

74/100

Risks:

Needs clearer deployment strategy

Recommendations:

Validate capacity assumptions with load tests and deployment evidence.

Cross Examination Results

No major contradictions detected across specialist agents.

Final Verdict

Blocking Issues

- No blocking issues recorded.

Recommended Fixes

1. Improve custom algorithms to increase innovation score
2. Enhance project's overall impact on the business or users

Deployment Roadmap

1. Refactor code for better modularity and maintainability
2. Implement additional security measures to improve security score

Deployment Decision

Verdict: `CONDITIONAL_APPROVE`
Overall Score: 73/100
Risk Level: MEDIUM